

# On Real-Time STM Concurrency Control for Embedded Software: Improved Time Bounds

We consider software transactional memory (STM) concurrency control for embedded multicore real-time software, and present a novel contention manager for resolving transactional conflicts, called FBLT. We upper bound transactional retries and task response times, and show when FBLT has better schedulability than previous contention managers and lock-free synchronization. Our implementation in the Rochester STM framework/real-time Linux reveals that FBLT yields shorter or comparable retry costs than competitors.

Categories and Subject Descriptors: C.3 [Special-Purpose and Application-based Systems]: Real-time and embedded systems

General Terms: Design, Experimentation, Measurement

Additional Key Words and Phrases: Software transactional memory (STM), real-time contention manager

## ACM Reference Format:

ACM Trans. Embedd. Comput. Syst. , , Article ( ), 25 pages.

DOI = 10.1145/0000000.0000000 <http://doi.acm.org/10.1145/0000000.0000000>

## 1. INTRODUCTION

Embedded systems sense physical processes and control their behavior, typically through feedback loops. Since physical processes are concurrent, computations that control them must also be concurrent, enabling them to process multiple streams of sensor input and control multiple actuators, all concurrently. Often, such computations need to concurrently read/write shared data objects. They must also process sensor input and react, while satisfying time constraints.

The de facto standard for concurrent programming is the threads abstraction, and the de facto synchronization abstraction is locks. Lock-based concurrency control has significant programmability, scalability, and composability challenges [Herlihy 2006]. Transactional memory (TM) is an alternative synchronization model for shared memory objects that promises to alleviate these difficulties. With TM, code that read/write shared objects is organized as *memory transactions*, which execute speculatively, while logging changes made to objects. Two transactions conflict if they access the same object and one access is a write. When that happens, a contention manager (CM) [Guer-raoui et al. 2005] resolves the conflict by aborting one and allowing the other to commit, yielding (the illusion of) atomicity. Aborted transactions are re-started, after rolling back the changes. In addition to a simple programming model, TM provides performance comparable to locking and lock-free approaches, especially for high contention and read-dominated workloads (see an example TM system's performance in [Saha et al. 2006]), and is composable [Harris et al. 2008]. TM has been proposed in hard-

---

New Paper, Not an Extension of a Conference Paper.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or [permissions@acm.org](mailto:permissions@acm.org).

© ACM 1539-9087/-ART \$15.00

DOI 10.1145/0000000.0000000 <http://doi.acm.org/10.1145/0000000.0000000>

ware, called HTM, and in software, called STM, with the usual tradeoffs: HTM has lesser overhead, but needs transactional support in hardware; STM is available on any hardware.

Given STM's programmability, scalability, and composability advantages, it is a compelling concurrency control technique also for multicore embedded real-time software. However, this requires bounding transactional retries, as real-time threads, which subsume transactions, must satisfy time constraints. Retry bounds under STM are dependent on the CM policy at hand (analogous to the way thread response time bounds are OS scheduler-dependent).

Past real-time CM research (Section 5) has proposed resolving transactional contention using dynamic and fixed priorities of parent threads, resulting in EDF-based CM (ECM) and RMS-based CM (RCM) [Fahmy and Ravindran 2011; El-Shambakey and Ravindran 2012b; 2012a], which are intended to be used with global EDF (G-EDF) and global RMS (G-RMS) multicore real-time schedulers [Davis and Burns 2011], respectively. In particular, [El-Shambakey and Ravindran 2012b] shows that ECM and RCM achieve higher schedulability – i.e., greater number of task sets meeting their time constraints – than lock-free synchronization only under some ranges for the maximum atomic section length. That range is significantly expanded with the LCM contention manager in [El-Shambakey and Ravindran 2012a], increasing the coverage of STM's timeliness superiority. ECM, RCM, and LCM suffer from transitive retry and cannot handle multiple objects per transaction efficiently.

These limitations are overcome with the PNF contention manager [El-Shambakey 2012]. However, PNF requires a-priori knowledge of all objects accessed by each transaction. This significantly limits programmability, and is incompatible with dynamic STM implementations [Herlihy et al. 2003]. Additionally, PNF is a centralized CM, which increases overheads and retry costs, and has a complex implementation.

We propose the First Bounded, Last Timestamp (or FBLT) contention manager (Section 5). In contrast to PNF, FBLT does not require a-priori knowledge of objects accessed by transactions. Moreover, FBLT allows each transaction to access multiple objects with shorter transitive retry cost than ECM, RCM and LCM. Additionally, FBLT is a decentralized CM and does not use locks in its implementation. Implementation of FBLT is also simpler than PNF.

We establish FBLT's retry and response time upper bounds under G-EDF and G-RMA schedulers (Section 6). We also identify the conditions under which FBLT's schedulability is better than ECM, RCM, G-EDF/LCM, G-RMA/LCM, PNF, and lock-free synchronization (Section 7).

We implement FBLT and competitor CM techniques in the Rochester STM framework [Marathe et al. ] and conduct experimental studies (Section 8). Our results reveal that FBLT has shorter retry cost than ECM, RCM and LCM. FBLT has no priori knowledge of accessed objects by transactions. So, retry cost under FBLT is a little higher than PNF.

Thus, the paper's contribution is the FBLT contention manager with superior timeliness properties. FBLT, thus allows programmers to reap STM's significant programmability and composability benefits for a broader range of multicore embedded real-time software than what was previously possible.

## 2. RELATED WORK

Transactional-like concurrency control without using locks, for real-time systems, has been previously studied in the context of non-blocking data structures (e.g., [Anderson et al. 1995]). Despite their numerous advantages over locks (e.g., deadlock-freedom), their programmability has remained a challenge. Past studies show that they are best suited for simple data structures where their retry cost is competitive to the cost of

lock-based synchronization [Brandenburg et al. 2008]. In contrast, STM is semantically simpler [Herlihy 2006], and is often the only viable lock-free solution for complex data structures (e.g., red/black tree) [Fahmy 2010] and nested critical sections [Saha et al. 2006].

STM concurrency control for real-time systems has been previously studied in [Manson et al. 2006; Fahmy et al. 2009; Sarni et al. 2009; Schoeberl et al. 2010; Fahmy 2010; Barros and Pinho 2011; El-Shambakey and Ravindran 2012b; 2012a; El-Shambakey 2012].

[Manson et al. 2006] proposes a restricted version of STM for uniprocessors. Uniprocessors do not need contention management. [Fahmy et al. 2009] bounds response times in distributed systems with STM synchronization. They consider Pfair scheduling, limit to small atomic regions with fixed size, and limit transaction execution to span at most two quanta. In contrast, we allow transaction lengths with arbitrary duration.

[Sarni et al. 2009] presents real-time scheduling of transactions and serializes transactions based on deadlines. However, the work does not bound retries and response times. In contrast, we establish such bounds. [Schoeberl et al. 2010] proposes real-time HTM. The work does not describe how transactional conflicts are resolved. Besides, the retry bound assumes that the worst case conflict between atomic sections of different tasks occurs when the sections are released at the same time. However, we show that this is not the worst case. We develop retry and response time upper bounds based on much worse conditions.

[Fahmy 2010] upper bounds retries and response times for ECM with G-EDF, and identify the tradeoffs with locking and lock-free protocols. Similar to [Schoeberl et al. 2010], [Fahmy 2010] also assumes that the worst case conflict between atomic sections occurs when the sections are released simultaneously. The ideas in [Fahmy 2010] are extended in [Barros and Pinho 2011], which presents three real-time CM designs. But no retry bounds or schedulability analysis techniques are presented for those CMs.

[El-Shambakey and Ravindran 2012b] presents the ECM and RCM contention managers, and upper bounds transactional retries and task response times under them. The work also identifies the conditions under which ECM and RCM are superior to locking and lock-free techniques. In particular, [El-Shambakey and Ravindran 2012b] shows that, STM's superiority holds only under some ranges for the maximum atomic section length. Moreover, [El-Shambakey and Ravindran 2012b] restricts transactions to access only one object. [El-Shambakey and Ravindran 2012a] presents length-based CM (LCM), and upper bounds transactional retries and response time for G-EDF/LCM and G-RMA/LCM. [El-Shambakey and Ravindran 2012a] compares (analytically and experimentally) between ECM, RCM and LCM, as well as lock-free and LCM. [El-Shambakey and Ravindran 2012a], as [El-Shambakey and Ravindran 2012b], restricts transactions to access only one object.

[El-Shambakey 2012] presents Priority CM with Negative value and First Access (PNF). PNF was designed to avoid transitive retry effect when each transaction accesses multiple objects. PNF also optimizes processor usage by reducing priority of retrying transactions below priorities of any real-time task. PNF requires prior knowledge of accessed objects by each transaction which is not always available. Besides, PNF is a centralized CM that uses locks in its implementation. Accordingly, reduction in retry cost is wasted by high overhead. [El-Shambakey 2012] upper bounds transactional retries and response time for G-EDF and G-RMA. [El-Shambakey 2012] compares (analytically and experimentally) between PNF and ECM, RCM, LCM and lock-free.

Our work builds upon [El-Shambakey and Ravindran 2012b; 2012a; El-Shambakey 2012], allows multiple objects per transaction with no prior knowledge about these

objects. We upper bound transactional retries and task response times. We identify the conditions for better schedulability for FBLT than other synchronization techniques.

### 3. PRELIMINARIES

We consider a multiprocessor system with  $m$  identical processors and  $n$  sporadic tasks  $\tau_1, \tau_2, \dots, \tau_n$ . The  $k^{th}$  instance (or job) of a task  $\tau_i$  is denoted  $\tau_i^k$ . Each task  $\tau_i$  is specified by its worst case execution time (WCET)  $c_i$ , its minimum period  $T_i$  between any two consecutive instances, and its relative deadline  $D_i$ , where  $D_i = T_i$ . Job  $\tau_i^j$  is released at time  $r_i^j$  and must finish no later than its absolute deadline  $d_i^j = r_i^j + D_i$ . Under a fixed priority scheduler such as G-RMA,  $p_i$  determines  $\tau_i$ 's (fixed) priority and it is constant for all instances of  $\tau_i$ . Under a dynamic priority scheduler such as G-EDF, a job  $\tau_i^j$ 's priority,  $p_i^j$ , differs from one instance to another. A task  $\tau_j$  may interfere with task  $\tau_i$  for a number of times during an interval  $L$ , and this number is denoted as  $G_{ij}(L)$ .

*Shared objects.* A task may need to read/write shared, in-memory data objects while it is executing any of its atomic sections (transactions), which are synchronized using STM. The set of atomic sections of task  $\tau_i$  is denoted  $s_i$ .  $s_i^k$  is the  $k^{th}$  atomic section of  $\tau_i$ . Each object,  $\theta$ , can be accessed by multiple tasks. The set of distinct objects accessed by  $\tau_i$  is  $\theta_i$  without repeating objects. The set of atomic sections used by  $\tau_i$  to access  $\theta$  is  $s_i(\theta)$ , and the sum of the lengths of those atomic sections is  $len(s_i(\theta))$ .  $s_i^k(\theta)$  is the  $k^{th}$  atomic section of  $\tau_i$  that accesses  $\theta$ .  $s_i^k$  can access one or more objects in  $\theta_i$ . So,  $s_i^k$  refers to the transaction itself, regardless of the objects accessed by the transaction. We denote the set of all accessed objects by  $s_i^k$  as  $\Theta_i^k$ . While  $s_i^k(\theta)$  implies that  $s_i^k$  accesses an object  $\theta \in \Theta_i^k$ ,  $s_i^k(\Theta)$  implies that  $s_i^k$  accesses a set of objects  $\Theta = \{\theta \in \Theta_i^k\}$ .  $s_i^k = s_i^k(\Theta)$  refers only once to  $s_i^k$ , regardless of the number of objects in  $\Theta$ . So,  $|s_i^k(\Theta)|_{\forall \theta \in \Theta} = 1$ .  $s_i^k(\theta)$  executes for a duration  $len(s_i^k(\theta))$ .  $len(s_i^k) = len(s_i^k(\theta)) = len(s_i^k(\Theta)) = len(s_i^k(\Theta_i^k))$ . The set of tasks sharing  $\theta$  with  $\tau_i$  is denoted  $\gamma_i(\theta)$ .

Atomic sections are non-nested (supporting nested STM is future work). The maximum-length atomic section in  $\tau_i$  that accesses  $\theta$  is denoted  $s_{i_{max}}(\theta)$ , while the maximum one among all tasks is  $s_{max}(\theta)$ , and the maximum one among tasks with priorities lower than that of  $\tau_i$  is  $s_{i_{max}}^l(\theta)$ .

*STM retry cost.* If two or more atomic sections conflict, the CM will commit one section and abort and retry the others, increasing the time to execute the aborted sections. The increased time that an atomic section  $s_i^p(\theta)$  will take to execute due to a conflict with another section  $s_j^k(\theta)$ , is denoted  $W_i^p(s_j^k(\theta))$ . If an atomic section,  $s_i^p$ , is already executing, and another atomic section  $s_j^k$  tries to access a shared object with  $s_i^p$ , then  $s_j^k$  is said to “interfere” or “conflict” with  $s_i^p$ . The job  $s_j^k$  is the “interfering job”, and the job  $s_i^p$  is the “interfered job”.

Due to *transitive retry* (introduced in Section 4.3), an atomic section  $s_i^k(\Theta_i^k)$  may retry due to another atomic section  $s_j^l(\Theta_j^l)$ , where  $\Theta_i^k \cap \Theta_j^l = \emptyset$ .  $\theta_i^*$  denotes the set of objects not accessed directly by atomic sections in  $\tau_i$ , but can cause transactions in  $\tau_i$  to retry due to transitive retry.  $\theta_i^{ex}(= \theta_i + \theta_i^*)$  is the set of all objects that can cause transactions in  $\tau_i$  to retry directly or through transitive retry.  $\gamma_i^*$  is the set of tasks that accesses objects in  $\theta_i^*$ .  $\gamma_i^{ex}(= \gamma_i + \gamma_i^*)$  is the set of all tasks that can directly or indirectly (through transitive retry) cause transactions in  $\tau_i$  to abort and retry.

The total time that a task  $\tau_i$ 's atomic sections have to retry over  $T_i$  is denoted  $RC(T_i)$ . The additional amount of time by which all interfering jobs of  $\tau_j$  increases the response time of any job of  $\tau_i$  during  $L$ , without considering retries due to atomic sections, is denoted  $W_{ij}(L)$ .

#### 4. MOTIVATION

To understand the need for *First Bounded, Last Timestamp (FBLT)* contention manager, we first give a brief introduction of previous real-time CMs, ECM [El-Shambakey and Ravindran 2012b], RCM [El-Shambakey and Ravindran 2012b], LCM [El-Shambakey and Ravindran 2012a], and PNF [El-Shambakey 2012].

##### 4.1. ECM and RCM

The *Earliest Deadline Contention Manager (ECM)* [El-Shambakey and Ravindran 2012b] is used with the G-EDF multicore real-time scheduler. ECM allows the transaction with the shortest absolute deadline to commit first. The other transactions retry and abort.

The *Rate Monotonic Contention Manager (RCM)* [El-Shambakey and Ravindran 2012b] is used with the G-RMA scheduler. RCM allows the transaction with the shortest period to commit first. ECM and RCM maintain semantic consistency with the underlying scheduler.

##### 4.2. LCM

For both ECM and RCM,  $s_i^k(\theta)$  can be totally repeated if  $s_j^l(\theta)$  — which belongs to a higher priority job  $\tau_j^b$  than  $\tau_i^a$  — conflicts with  $s_i^k(\theta)$  at the end of its execution, while  $s_i^k(\theta)$  is just about to commit. The *Length-based Contention Manager (LCM)* [El-Shambakey and Ravindran 2012a], shown in Algorithm 1, uses the remaining length of  $s_i^k(\theta)$  when it is interfered, as well as  $len(s_j^l(\theta))$ , to decide which transaction must be aborted. If  $p_i^k$  is greater than  $p_j^l$ , then  $s_i^k(\theta)$  is committed, because it belongs to a higher priority job, and it started before  $s_j^l(\theta)$  (step 2). Otherwise,  $c_{ij}^{kl}$  is calculated (step 4) to determine whether it is worth aborting  $s_i^k(\theta)$  in favor of  $s_j^l(\theta)$ , because  $len(s_j^l(\theta))$  is relatively small compared to the remaining execution length of  $s_i^k(\theta)$ . [El-Shambakey and Ravindran 2012a] assumes that:

$$c_{ij}^{kl} = len(s_j^l(\theta)) / len(s_i^k(\theta)) \quad (1)$$

where  $c_{ij}^{kl} \in ]0, \infty[$ , to cover all possible lengths of  $s_j^l(\theta)$ .

Thus, LCM's key idea is to reduce the opportunity for the abort of  $s_i^k(\theta)$  if it is close to committing when interfered and  $len(s_j^l(\theta))$  is large. This abort opportunity is increasingly reduced as  $s_i^k(\theta)$  gets closer to the end of its execution, or  $len(s_j^l(\theta))$  gets larger. On the other hand, as  $s_i^k(\theta)$  is interfered early, or  $len(s_j^l(\theta))$  is small compared to  $s_i^k(\theta)$ 's remaining length, the abort opportunity is increased even if  $s_i^k(\theta)$  is close to the end of its execution. To decide whether  $s_i^k(\theta)$  must be aborted or not, a threshold value  $\psi \in [0, 1]$  that determines  $\alpha_{ij}^{kl}$  (step 5) is used, where  $\alpha_{ij}^{kl}$  is the maximum percentage of  $len(s_i^k(\theta))$  below which  $s_j^l(\theta)$  is allowed to abort  $s_i^k(\theta)$ . Thus, if the already executed part of  $s_i^k(\theta)$  — when  $s_j^l(\theta)$  interferes with  $s_i^k(\theta)$  — does not exceed  $\alpha_{ij}^{kl} len(s_i^k(\theta))$ , then  $s_i^k(\theta)$  is aborted (step 8). Otherwise,  $s_j^l(\theta)$  is aborted (step 10).

LCM reduces the retry cost of a single transaction  $s_i^k(\theta)$  due to another transaction  $s_j^l(\theta)$  from  $2 \cdot s_{max}$  (in case of ECM and RCM) to  $(1 + \alpha_{max}) \cdot s_{max}$ , where  $s_{max}$  is the maximum length transaction among all tasks, and  $\alpha_{max}$  is the maximum *alpha* for any transaction. On the other hand, LCM suffers from bounded priority inversion because a higher priority transaction can be blocked by a lower priority one. Additionally, LCM is not a centralized CM, which means that, upon a conflict, each transaction must decide whether it must commit or abort.

**ALGORITHM 1: The LCM Algorithm****Data:**  $s_i^k(\theta) \rightarrow$  interfered atomic section. $s_j^l(\theta) \rightarrow$  interfering atomic section. $\psi \rightarrow$  predefined threshold  $\in [0, 1]$ . $\delta_i^k(\theta) \rightarrow$  remaining execution length of  $s_i^k(\theta)$ **Result:** which atomic section of  $s_i^k(\theta)$  or  $s_j^l(\theta)$  aborts

```

1 if  $p_i^k > p_j^l$  then
2   |  $s_j^l(\theta)$  aborts;
3 else
4   |  $c_{ij}^{kl} = \text{len}(s_j^l(\theta)) / \text{len}(s_i^k(\theta));$ 
5   |  $\alpha_{ij}^{kl} = \ln(\psi) / (\ln(\psi) - c_{ij}^{kl});$ 
6   |  $\alpha = (\text{len}(s_i^k(\theta)) - \delta_i^k(\theta)) / \text{len}(s_i^k(\theta));$ 
7   | if  $\alpha \leq \alpha_{ij}^{kl}$  then
8     |  $s_i^k(\theta)$  aborts;
9   | else
10    |  $s_j^l(\theta)$  aborts;
11  | end
12 end

```

**4.3. PNF**

ECM, RCM, and LCM suffer from *transitive retry*. Transitive retry is illustrated by the following example:

Consider three atomic sections  $s_1^x$ ,  $s_2^y$ , and  $s_3^z$  belonging to jobs  $\tau_1^x$ ,  $\tau_2^y$ , and  $\tau_3^z$ , with priorities  $p_3^z > p_2^y > p_1^x$ , respectively. Assume that  $s_1^x$  and  $s_2^y$  share objects, and  $s_2^y$  and  $s_3^z$  share objects.  $s_1^x$  and  $s_3^z$  do not share objects. Now,  $s_3^z$  can cause  $s_2^y$  to retry, which in turn will cause  $s_1^x$  to retry. This means that  $s_1^x$  will retry transitively because of  $s_3^z$ , which will increase the retry cost of  $s_1^x$ .

Now, consider another atomic section  $s_4^f$  with a priority higher than that of  $s_3^z$ . Suppose  $s_4^f$  shares objects only with  $s_3^z$ . Thus,  $s_4^f$  can cause  $s_3^z$  to retry, which in turn will cause  $s_2^y$  to retry, and finally,  $s_1^x$  to retry. Thus, transitive retry will move from  $s_4^f$  to  $s_1^x$ , increasing the retry cost of  $s_1^x$ . The situation gets worse as more higher priority tasks are added, where each task shares objects with its immediate lower priority task.  $\tau_3^z$  may have atomic sections that share objects with  $\tau_1^x$ , but this will not prevent the effect of transitive retry due to  $s_1^x$ .

Therefore, the analysis in [El-Shambakey and Ravindran 2012b] and [El-Shambakey and Ravindran 2012a] extend the set of objects that can cause an atomic section of a lower priority job to retry. This is done by initializing the set of conflicting objects,  $\gamma_i$ , to all objects accessed by all transactions of  $\tau_i$ . We then cycle through all transactions belonging to all other higher priority tasks. Each transaction  $s_j^l$  that accesses at least one of the objects in  $\gamma_i$  adds all other objects accessed by  $s_j^l$  to  $\gamma_i$ . The loop over all higher priority tasks is repeated, each time with the new  $\gamma_i$ , until there are no more transactions accessing any object in  $\gamma_i$ . The final set of objects (tasks) that can cause transactions in  $\tau_i$  to retry is  $\theta_i^{ex}(\gamma_i^{ex})$ , respectively<sup>1</sup>.

The *Priority contention manager with Negative value and First access* (PNF) [El-Shambakey 2012] is designed to avoid transitive retry. ECM, RCM, and LCM suffer from transitive retry. PNF avoids transitive retry by concurrently executing at most  $m$  non-conflicting transactions together. These executing transactions are non-

<sup>1</sup>However, note that, this solution may over-extend the set of conflicting objects, and may even contain all objects accessed by all tasks.

preemptive. Thus, executing transactions cannot be aborted due to direct or indirect conflict with other transactions.

However, with PNF, all objects accessed by each transaction must be known a-priori. Therefore, this is not suitable with dynamic STM implementations [Herlihy et al. 2003]. Additionally, PNF is implemented in [El-Shambakey 2012] as a centralized CM that uses locks. This increases overhead.

#### 4.4. Case for FBLT

Thus, it is desirable to have a CM with the following goals:

- (1) reduce the retry cost of each transaction  $s_i^k$  due to another transaction  $s_j^l$ , just as LCM does compared to ECM and RCM.
- (2) avoid or bound the effect of transitive retry, similar to PNF, without prior knowledge of accessed objects by each transaction, enabling dynamic STM.
- (3) decentralized design and avoid the use of locks, thereby reducing overhead.

We propose the *First Bounded, Last Timestamp contention manager* (or (FBLT)). FBLT achieves these goals by bounding the number of times each transaction  $s_i^k$  is aborted due to other transactions to at most  $\delta_i^k$ .  $\delta_i^k$  includes the number of aborts due to direct conflict with other transactions, as well as transitive retry (goal 2). If a transaction  $s_i^k$  reaches its  $\delta_i^k$ , it is added to an  $m\_set$  in FIFO order. In the  $m\_set$ ,  $s_i^k$  executes non-preemptively. If transactions in the  $m\_set$  conflict together, they use their FIFO order in the  $m\_set$  to resolve the conflict.  $s_i^k$  can still abort after it becomes a non-preemptive transaction due to other non-preemptive transactions. The number of aborts for any non-preemptive transaction is bounded by  $m - 1$ , where  $m$  is the number of processors, as will be shown in Section 6.

Thus, the key idea behind FBLT is to use a suitable  $\delta_i^k$  for each  $s_i^k$  before it becomes a non-preemptive transaction. The choice of  $\delta_i^k$  should make the total retry cost (and thus, the schedulability) of any job  $\tau_i^x$  under FBLT comparable to the retry cost under ECM, RCM, LCM, and PNF. (In Section 7, we show the suitable  $\delta_i^k$  for each  $s_i^k$  to have equal or better schedulability than other CMs.) Preemptive transactions resolve their conflicts using LCM. Thus, FBLT defaults to LCM if abort bounds have not been violated (goal 1). Each non-preemptive transaction  $s_i^k$  uses the time it joined the  $m\_set$  to resolve conflicts with other non-preemptive transactions. Therefore, FBLT does not have to use locks and is decentralized (goal 3).

### 5. THE FBLT CONTENTION MANAGER

Algorithm 2 illustrates FBLT. Each transaction  $s_i^k$  can be aborted during  $T_i$  for at most  $\delta_i^k$  times.  $\eta_i^k$  records the number of times  $s_i^k$  has already been aborted up to now. If  $s_i^k$  and  $s_j^l$  have not joined the  $m\_set$  yet, then they are preemptive transactions. Preemptive transactions resolve conflicts using Algorithm 1 (step 2). Thus, FBLT defaults to LCM when no transaction reaches its  $\delta$ . If only one of the transactions is in the  $m\_set$ , then the non-preemptive transaction (the one in  $m\_set$ ) aborts the other one (steps 15 to 26).  $\eta_i^k$  is incremented each time  $s_i^k$  is aborted as long as  $\eta_i^k < \delta_i^k$  (steps 5 and 18). Otherwise,  $s_i^k$  is added to the  $m\_set$  and its priority is increased to  $m\_prio$  (steps 7 to 9 and 20 to 22). When the priority of  $s_i^k$  is increased to  $m\_prio$ ,  $s_i^k$  becomes a non-preemptive transaction. Non-preemptive transactions cannot be aborted by other preemptive transactions, nor by any other real-time job. The  $m\_set$  can hold at most  $m$  concurrent transactions because there are  $m$  processors in the system.  $r(s_i^k)$  records the time  $s_i^k$  joined the  $m\_set$  (steps 8 and 21). When non-preemptive transactions conflict together (step 27), the transaction with the smaller  $r()$  commits first (steps 29 and 31). Thus, non-preemptive transactions are executed in FIFO order of the  $m\_set$ .

**ALGORITHM 2: The FBLT Algorithm****Data:**  $s_i^k$ : interfered transaction; $s_j^l$ : interfering transactions; $\delta_i^k$ : the maximum number of times  $s_i^k$  can be aborted during  $T_i$ ; $\eta_i^k$ : number of times  $s_i^k$  has already been aborted up to now; $m\_set$ : contains at most  $m$  non-preemptive transactions.  $m$  is number of processors; $m\_prio$ : priority of any transaction in  $m\_set$ .  $m\_prio$  is higher than any priority of any real-time task; $r(s_i^k)$ : time point at which  $s_i^k$  joined  $m\_set$ ;**Result:** atomic sections that will abort

```

1 if  $s_i^k, s_j^l \notin m\_set$  then
2   Apply Algorithm 1 (default to LCM);
3   if  $s_i^k$  is aborted then
4     if  $\eta_i^k < \delta_i^k$  then
5       Increment  $\eta_i^k$  by 1;
6     else
7       Add  $s_i^k$  to  $m\_set$ ;
8       Record  $r(s_i^k)$ ;
9       Increase priority of  $s_i^k$  to  $m\_prio$ ;
10    end
11  else
12    Swap  $s_i^k$  and  $s_j^l$ ;
13    Go to Step 3;
14  end
15 else if  $s_j^l \in m\_set, s_i^k \notin m\_set$  then
16   Abort  $s_i^k$ ;
17   if  $\eta_i^k < \delta_i^k$  then
18     Increment  $\eta_i^k$  by 1;
19   else
20     Add  $s_i^k$  to  $m\_set$ ;
21     Record  $r(s_i^k)$ ;
22     Increase priority of  $s_i^k$  to  $m\_prio$ ;
23   end
24 else if  $s_i^k \in m\_set, s_j^l \notin m\_set$  then
25   Swap  $s_i^k$  and  $s_j^l$ ;
26   Go to Step 15;
27 else
28   if  $r(s_i^k) < r(s_j^l)$  then
29     Abort  $s_j^l$ ;
30   else
31     Abort  $s_i^k$ ;
32   end
33 end

```

**5.1. Illustrative Example**

We now illustrate FBLT's behavior with the following example:

- (1) Transaction  $s_i^k(\theta_1, \theta_2)$  is released while  $m\_set = \emptyset$ .  $\eta_i^k = 0$  and  $\delta_i^k = 3$ .
- (2) Transaction  $s_a^b(\theta_2)$  is released while  $s_i^k(\theta_1, \theta_2)$  is running.  $p_a^b > p_i^k$  and  $\eta_i^k < \delta_i^k$ . Applying LCM,  $s_i^k(\theta_1, \theta_2)$  is aborted in favor of  $s_a^b$  and  $\eta_i^k$  is incremented to 1.
- (3)  $s_a^b(\theta_2)$  commits.  $s_i^k(\theta_1, \theta_2)$  runs again. Transaction  $s_c^d(\theta_2)$  is released while  $s_i^k(\theta_1, \theta_2)$  is running.  $p_c^d > p_i^k$ . Applying LCM,  $s_i^k(\theta_1, \theta_2)$  is aborted again in favor of  $s_c^d(\theta_2)$ .  $\eta_i^k$  is incremented to 2.



- (4)  $s_e^d(\theta_2)$  commits.  $s_e^f(\theta_2, \theta_3)$  is released.  $p_e^f > p_i^k$  and  $\eta_e^f = 2$ .  $s_i^k(\theta_1, \theta_2)$  is aborted in favor of  $s_e^f(\theta_2, \theta_3)$  and  $\eta_i^k$  is incremented to 3.
- (5)  $s_j^l(\theta_3)$  is released.  $p_j^l > p_e^f$ .  $s_e^f(\theta_2, \theta_3)$  is aborted in favor of  $s_j^l(\theta_3)$  and  $\eta_e^f$  is incremented to 1.
- (6)  $s_i^k(\theta_1, \theta_2)$  and  $s_e^f(\theta_2, \theta_3)$  are compared again.  $\eta_i^k = \delta_i^k$ ,  $s_i^k(\theta_1, \theta_2)$  is added to  $m\_set$ .  $m\_set = \{s_i^k(\theta_1, \theta_2)\}$ .  $s_i^k(\theta_1, \theta_2)$  becomes a non-preemptive transaction. As  $s_e^f(\theta_2, \theta_3)$  is a preemptive transaction,  $s_e^f(\theta_2, \theta_3)$  is aborted in favor of  $s_i^k(\theta_1, \theta_2)$ , despite  $p_e^f$  being greater than the original priority of  $s_i^k(\theta_1, \theta_2)$ .  $\eta_e^f$  is incremented to 2.
- (7)  $s_j^l(\theta_3)$  commits but  $s_g^h(\theta_3)$  is released.  $p_g^h > p_e^f$  but  $\eta_e^f = \delta_e^f$ . So,  $s_e^f(\theta_2, \theta_3)$  becomes a non-preemptive transaction.  $m\_set = \{s_i^k(\theta_1, \theta_2), s_g^h(\theta_2, \theta_3)\}$ .
- (8)  $s_i^k(\theta_1, \theta_2)$  and  $s_g^h(\theta_2, \theta_3)$  are now non-preemptive transactions.  $s_i^k(\theta_1, \theta_2)$  and  $s_g^h(\theta_2, \theta_3)$  still conflict together. So, they are executed according to their addition order to the  $m\_set$ . So,  $s_i^k(\theta_1, \theta_2)$  commits first, followed  $s_g^h(\theta_2, \theta_3)$ .
- (9)  $s_g^h(\theta_3)$  will continue to abort and retry in favor of  $s_e^f(\theta_2, \theta_3)$  until  $s_e^f(\theta_2, \theta_3)$  commits or  $\eta_g^h = \delta_g^h$ . Even if  $s_g^h(\theta_3)$  joined the  $m\_set$ ,  $s_g^h(\theta_3)$  will still abort and retry in favor of  $s_e^f(\theta_2, \theta_3)$ , because  $s_e^f(\theta_2, \theta_3)$  joined the  $m\_set$  earlier than  $s_g^h(\theta_3)$ .

It is seen from steps 2 to 6 that  $s_i^k(\theta_1, \theta_2)$  can be aborted due to direct conflict with other transactions, or due to transitive retry. Irrespective of the reason for the conflict, once a transaction has reached its maximum allowed  $\delta$ , the transaction becomes a non-preemptive one (steps 6 and 7). Non-preemptive transactions have higher priority than other preemptive transactions (steps 6 and 7). Non-preemptive transactions execute in their arrival order to the  $m\_set$ .

## 6. RETRY COST AND RESPONSE TIME BOUNDS

We now derive an upper bound on the retry cost of any job  $\tau_i^x$  under FBLT during an interval  $L \leq T_i$ . Since all tasks are sporadic (i.e., each task  $\tau_i$  has a minimum period  $T_i$ ),  $T_i$  is the maximum study interval for each task  $\tau_i$ .

**CLAIM 1.** *The total retry cost for any job  $\tau_i^x$  under FBLT due to 1) conflicts between its transactions and transactions of other jobs during an interval  $L \leq T_i$  and 2) release of higher priority jobs is upper bounded by:*

$$RC_{to}(L) \leq \sum_{\forall s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{\forall s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right) + RC_{re}(L) \quad (2)$$

where  $\chi_i^k$  is the set of at most  $m - 1$  maximum length transactions conflicting directly or indirectly (through transitive retry) with  $s_i^k$ . Each transaction  $s_{iz}^k \in \chi_i^k$  belongs to a distinct task  $\tau_j$ .  $RC_{re}(L)$  is the retry cost resulting from the release of higher priority jobs which preempt  $\tau_i^x$ .  $RC_{re}(L)$  is calculated by (6.8) in [El-Shambakey 2012] for G-EDF, and (6.10) in [El-Shambakey 2012] for G-RMA schedulers.

**PROOF.** By the definition of FBLT,  $s_i^k \in \tau_i^x$  can be aborted a maximum of  $\delta_i^k$  times before  $s_i^k$  joins the  $m\_set$ . Before joining the  $m\_set$ ,  $s_i^k$  can be aborted due to higher priority transactions, or transactions in the  $m\_set$ . The original priority of transactions in the  $m\_set$  can be higher or lower than  $p_i^x$ . Thus, the maximum time  $s_i^k$  is aborted before joining the  $m\_set$  occurs if  $s_i^k$  is aborted for  $\delta_i^k$  times.

Transactions preceding  $s_i^k$  in the  $m\_set$  can conflict directly with  $s_i^k$ , or indirectly through transitive retry. The worst case scenario for  $s_i^k$  after joining the  $m\_set$  occurs if  $s_i^k$  is preceded by  $m - 1$  maximum length conflicting transactions. Hence, in the worst

case,  $s_i^k$  has to wait for the previous  $m - 1$  transactions to commit first. The priority of  $s_i^k$  after joining the  $m\_set$  is higher than any real-time job. Therefore,  $s_i^k$  is not aborted by any job. If  $s_i^k$  has not joined the  $m\_set$  yet, and a higher priority job  $\tau_j^y$  is released while  $s_i^k$  is running, then  $s_i^k$  may be aborted if  $\tau_j^y$  has conflicting transactions with  $s_i^k$ .  $\tau_j^y$  causes only one abort in  $\tau_i^x$  because  $\tau_j^y$  preempts  $\tau_i^x$  only once. If  $s_i^k$  has already joined the  $m\_set$ , then  $s_i^k$  cannot be aborted by the release of higher priority jobs. Thus, the maximum number of times transactions in  $\tau_i^x$  can be aborted due to the release of higher priority jobs is less than or equal to the number of interfering higher priority jobs to  $\tau_i^x$ . Claim follows.  $\square$

CLAIM 2. *Under FBLT, the blocking time of a job  $\tau_i^x$  due to lower priority jobs is upper bounded by:*

$$D(\tau_i^x) = \min \left( \max_1^m (s_{j_{max}, \forall \tau_j^l, p_j^l < p_i^x}) \right) \quad (3)$$

where  $s_{j_{max}}$  is the maximum length transaction in any job  $\tau_j^l$  with original priority lower than  $p_i^x$ . The right hand side of (6) is the minimum of the  $m$  maximum transactional lengths in all jobs with lower priority than  $\tau_i^x$ .

PROOF.  $\tau_i^x$  is blocked when it is initially released and all processors are busy with lower priority jobs with non-preemptive transactions. Although  $\tau_i^x$  can be preempted by higher priority jobs,  $\tau_i^x$  cannot be blocked after it is released. If  $\tau_i^x$  is preempted by a higher priority job  $\tau_j^y$ , then, when  $\tau_j^y$  finishes execution, the underlying scheduler will not choose a lower priority job than  $\tau_i^x$  before  $\tau_i^x$ . So, after  $\tau_i^x$  is released, there is no chance for any transaction  $s_u^v$  belonging to a lower priority job than  $\tau_i^x$  to run before  $\tau_i^x$ . Thus,  $s_u^v$  cannot join the  $m\_set$  before  $\tau_i^x$  finishes. Consequently, the worst case blocking time for  $\tau_i^x$  occurs when the maximum length  $m$  transactions in lower priority jobs than  $\tau_i^x$  are executing non-preemptively. After the minimum length transaction in the  $m\_set$  finishes, the underlying scheduler will choose  $\tau_i^x$  or a higher priority job to run. Claim follows.  $\square$

CLAIM 3. *The response time of any job  $\tau_i^x$  during an interval  $L \leq T_i$  under FBLT is upper bounded by:*

$$R_i^{up} = c_i + RC_{to}(L) + D(\tau_i^x) + \left\lceil \frac{1}{m} \sum_{\forall j \neq i} W_{ij}(R_i^{up}) \right\rceil \quad (4)$$

where  $RC_{to}(L)$  is calculated by (5),  $D(\tau_i^x)$  is calculated by (6), and  $W_{ij}(R_i^{up})$  is calculated by (11) in [El-Shambakey and Ravindran 2012b] for G-EDF, and (17) in [El-Shambakey and Ravindran 2012b] for G-RMA schedulers. (11) and (17) in [El-Shambakey and Ravindran 2012b] inflates  $c_j$  of any job  $\tau_j^y \neq \tau_i^x$ ,  $p_j^y > p_i^x$  by the retry cost of transactions in  $\tau_j^y$ .

PROOF. The response time of a job is calculated directly from FBLT's behavior. The response time of any job  $\tau_i^x$  is the sum of its worst case execution time  $c_i$ , plus the retry cost of transactions in  $\tau_i^x$  ( $RC_{to}(L)$ ), plus the blocking time of  $\tau_i^x$  ( $D(\tau_i^x)$ ), and the workload interference of higher priority jobs. The workload interference of higher priority jobs scheduled by G-EDF is calculated by (11) in [El-Shambakey and Ravindran 2012b], and by (17) in [El-Shambakey and Ravindran 2012b] for G-RMA. Claim follows.  $\square$

We now derive an upper bound on the retry cost of any job  $\tau_i^x$  under FBLT during an interval  $L \leq T_i$ . Since all tasks are sporadic (i.e., each task  $\tau_i$  has a minimum period  $T_i$ ),  $T_i$  is the maximum study interval for each task  $\tau_i$ .

CLAIM 4.

*The total retry cost for any job  $\tau_i^x$  under FBLT due to: 1) conflicts between its transactions and transactions of other jobs during an interval  $L \leq T_i$ . 2) release of higher priority jobs, is upper bounded by:*

$$RC_{to}(L) \leq \sum_{\forall s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{\forall s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right) + RC_{re}(L) \quad (5)$$

where  $\chi_i^k$  is the set of at most  $m - 1$  maximum length transactions conflicting directly or indirectly (through transitive retry) with  $s_i^k$ . Each transaction  $s_{iz}^k \in \chi_i^k$  belongs to a distinct task  $\tau_j$ .  $RC_{re}(L)$  is the retry cost resulting from release of higher priority jobs which preempt  $\tau_i^x$ .  $RC_{re}(L)$  is calculated by (6.8) in [El-Shambakey 2012] for G-EDF, and (6.10) in [El-Shambakey 2012] for G-RMA.

PROOF.

By definition of FBLT,  $s_i^k \in \tau_i^x$  can be aborted at maximum  $\delta_i^k$  times before  $s_i^k$  joins  $m\_set$ . Before joining  $m\_set$ ,  $s_i^k$  can be aborted due to higher priority transactions, or transactions in the  $m\_set$ . Original priority of transactions in  $m\_set$  can be of higher or lower priority than  $p_i^x$ . Thus, the maximum time  $s_i^k$  is aborted before joining  $m\_set$  occurs if  $s_i^k$  is aborted for  $\delta_i^k$  times. Transactions preceding  $s_i^k$  in  $m\_set$  can conflict directly with  $s_i^k$ , or indirectly through transitive retry. The worst case scenario for  $s_i^k$  after joining  $m\_set$  occurs if  $s_i^k$  is preceded by  $m - 1$  maximum length conflicting transactions. Hence, in worst case,  $s_i^k$  has to wait for the previous  $m - 1$  transactions to commit first. Priority of  $s_i^k$  after joining  $m\_set$  is higher than any real-time job. So,  $s_i^k$  is not aborted by any job. If  $s_i^k$  has not joined  $m\_set$  yet, and a higher priority job  $\tau_j^y$  is released while  $s_i^k$  is running, then  $s_i^k$  may be aborted if  $\tau_j^y$  has conflicting transactions with  $s_i^k$ .  $\tau_j^y$  causes only one abort in  $\tau_i^x$  because  $\tau_j^y$  preempts  $\tau_i^x$  only once. If  $s_i^k$  has already joined  $m\_set$ , then  $s_i^k$  cannot be aborted by release of higher priority jobs. So, the maximum number of abort times to transactions in  $\tau_i^x$  due to release of higher priority jobs is less or equal to number of interfering higher priority jobs to  $\tau_i^x$ . Claim follows.

□

CLAIM 5.

*The blocking time for a job  $\tau_i^x$  due to lower priority jobs is upper bounded by:*

$$D(\tau_i^x) = \min \left( \max_1^m (s_{j_{max}, \forall \tau_j^l, p_j^l < p_i^x}) \right) \quad (6)$$

where  $s_{j_{max}}$  is the maximum length transaction in any job  $\tau_j^l$  with original priority lower than  $p_i^x$ . The right hand side of (6) is the minimum of the  $m$  maximum transactional lengths in all jobs with lower priority than  $\tau_i^x$ .

PROOF.

$\tau_i^x$  is blocked when it is initially released and all processors are busy with lower priority jobs with non-preemptive transactions. Although  $\tau_i^x$  can be preempted by higher priority jobs,  $\tau_i^x$  cannot be blocked after it is released. If  $\tau_i^x$  is preempted by a higher

priority job  $\tau_j^y$ , then  $\tau_j^y$  finishes execution, the underlying scheduler will not choose a lower priority job than  $\tau_i^x$  before  $\tau_i^x$ . So, after  $\tau_i^x$  is released, there is no chance for any transaction  $s_u^v$  belonging to a lower priority job than  $\tau_i^x$  to run before  $\tau_i^x$ . Thus,  $s_u^v$  cannot join  $m\_set$  before  $\tau_i^x$  finishes. Consequently, the worst case blocking time for  $\tau_i^x$  occurs when the maximum length  $m$  transactions in lower priority jobs than  $\tau_i^x$  are executing non-preemptively. After the minimum length transaction in the  $m\_set$  finishes, the underlying scheduler will choose  $\tau_i^x$  or a higher priority job to run. Claim follows.

□

#### CLAIM 6.

*Response time of any job  $\tau_i^x$  during an interval  $L \leq T_i$  under FBLT is upper bounded by*

$$R_i^{up} = c_i + RC_{to}(L) + D(\tau_i^x) + \left\lfloor \frac{1}{m} \sum_{\forall j \neq i} W_{ij}(R_i^{up}) \right\rfloor \quad (7)$$

where  $RC_{to}(L)$  is calculated by (5),  $D(\tau_i^x)$  is calculated by (6), and  $W_{ij}(R_i^{up})$  is calculated by (11) in [El-Shambakey and Ravindran 2012b] for G-EDF, and (17) in [El-Shambakey and Ravindran 2012b] for G-RMA. (11) and (17) in [El-Shambakey and Ravindran 2012b] inflates  $c_j$  of any job  $\tau_j^y \neq \tau_i^x$ ,  $p_j^y > p_i^x$  by retry cost of transactions in  $\tau_j^y$ .

#### PROOF.

Response time of any job  $\tau_i^x$  is calculated directly from FBLT's behaviour. Response time of any job  $\tau_i^x$  is the sum of its worst case execution time  $c_i$ , plus retry cost of transactions in  $\tau_i^x$  ( $RC_{to}(L)$ ), plus blocking time of  $\tau_i^x$  ( $D(\tau_i^x)$ ), and the workload interference of higher priority jobs. Workload interference of higher priority jobs scheduled by G-EDF is calculated by (11) in [El-Shambakey and Ravindran 2012b], and by (17) in [El-Shambakey and Ravindran 2012b] for G-RMA. Claim follows.

□

## 7. SCHEDULABILITY COMPARISON

We now (formally) compare the schedulability of G-EDF (G-RMA) with FBLT against ECM, RCM, LCM, PNF, and lock-free synchronization [El-Shambakey and Ravindran 2012b; 2012a; Devi et al. 2006; El-Shambakey 2012]. Such a comparison will reveal when FBLT outperforms the others. Toward this, we compare the total utilization under G-EDF (G-RMA)/FBLT with that under the other synchronization methods. In this comparison, we use the inflated execution time of the task, which is the sum of the worst-case execution time of the task and its retry cost, in the utilization calculation of the task.

Note that, for a job  $\tau_i^x$ , no processor is available during its blocking time. Since each processor is busy with some job other than  $\tau_i^x$ ,  $D(\tau_i^x)$  is not added to the inflated execution time of  $\tau_i^x$ . Hence,  $D(\tau_i^x)$  is not added to the utilization calculation of  $\tau_i^x$ .

Let  $RC_A(T_i)$  and  $RC_B(T_i)$  denote the retry cost of a job  $\tau_i^x$  during  $T_i$  using the synchronization method  $A$  and synchronization method  $B$ , respectively. Now, schedulability of  $A$  is comparable to  $B$  if:

$$\sum_{\forall \tau_i} \frac{c_i + RC_A(T_i)}{T_i} \leq \sum_{\forall \tau_i} \frac{c_i + RC_B(T_i)}{T_i}$$

$$\sum_{\forall \tau_i} \frac{RC_A(T_i)}{T_i} \leq \sum_{\forall \tau_i} \frac{RC_B(T_i)}{T_i} \quad (8)$$

### 7.1. FBLT vs. ECM

**CLAIM 7.** *The schedulability of FBLT is equal to or better than ECM's when the maximum abort number of any preemptive transaction  $s_i^k$  is less than or equal to the number of transactions directly conflicting with  $s_i^k$  in all other jobs with higher priority than  $\tau_i$ 's current job.*

**PROOF.**

By substituting  $RC_A(T_i)$  and  $RC_B(T_i)$  in (8) with (5) and (6.7) in [El-Shambakey 2012], respectively, we get:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{\forall s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right) + RC_{re}(T_i)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \sum_{\forall \tau_j \in \gamma_i^{ex}} \sum_{\theta \in \theta_i^{ex}} \left( \left\lceil \frac{T_i}{T_j} \right\rceil \sum_{\forall s_j^h(\theta)} \text{len}(\bar{s}_j^h(\theta) + s_{max}^j(\theta)) \right) \right) + RC_{re}(T_i)}{T_i} \end{aligned} \quad (9)$$

Let  $\theta_i^{ex} = \theta_i + \theta_i^*$ , where  $\theta_i^*$  is the set of objects not accessed directly by  $\tau_i$  but can cause transactions in  $\tau_i$  to retry due to transitive retry. Let  $\gamma_i^{ex} = \gamma_i + \gamma_i^*$ , where  $\gamma_i^*$  is the set of tasks that access objects in  $\theta_i^*$ .  $\bar{s}_j^h(\theta)$  can access multiple objects, so  $s_{max}^j(\theta)$  is the maximum length transaction conflicting with  $\bar{s}_j^h(\theta)$ .  $\bar{s}_j^h(\theta)$  is included only once for all  $\theta \in \Theta_j^h$ . Each  $\theta \in \theta_i^{ex}$  has its own  $s_{max}^j(\theta)$ . But  $s_i^h$  can access multiple objects, denoted as  $\Theta_j^h$ . So,  $s_{max}^j(\theta)$  is replaced by  $s_{max}^j(\Theta_j^h)$ , where  $s_{max}^j(\Theta_j^h) = \max\{s_{max}^j(\theta), \forall \theta \in \Theta_j^h\}$ .  $s_{max}^j(\Theta_j^h)$  is included once for each  $\theta \in \theta_i$ .

Each job  $\tau_i^x$  has the same interference pattern from higher priority jobs,  $\tau_j^h$ , under FBLT and ECM. Hence,  $RC_{re}(T_i)$  for  $\tau_i^x$  is the same under FBLT and ECM. Consequently, (9) becomes:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{\forall s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \sum_{\forall \tau_j \in \gamma_i^{ex}} \sum_{\forall s_j^h(\Theta_j^h), \Theta_j^h \in \theta_i^{ex}} \left( \left\lceil \frac{T_i}{T_j} \right\rceil \text{len}(\bar{s}_j^h(\Theta_j^h) + s_{max}^j(\Theta_j^h)) \right) \right)}{T_i} \end{aligned} \quad (10)$$

Although different  $s_i^k$ s can have common conflicting transactions  $\bar{s}_j^h$ , no more than one  $s_i^k$  can be preceded by the same  $\bar{s}_j^h$  in the  $m\_set$ . This happens because transactions in the  $m\_set$  are non-preemptive. The original priority of transactions preceding  $s_i^k$  in the  $m\_set$  can be lower or higher than the original priority of  $s_i^k$ . Since under G-EDF,  $\tau_j$  can have at least one job of higher priority than  $\tau_i^x$ ,  $\left\lceil \frac{T_i}{T_j} \right\rceil \geq 1$ . Thus, each one of the  $s_{iz}^k$  term in the left hand side of (10) is included in one of the  $\bar{s}_j^h(\theta)$  term in the right hand side of (10). Now, (10) holds if:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{\forall s_i^k \in s_i} \delta_i^k \text{len}(s_i^k)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\sum_{\forall \tau_j \in \gamma_i^{ex}} \sum_{\forall s_j^h(\Theta_j^h), \Theta_j^h \in \theta_i^{ex}} \left( \left\lceil \frac{T_i}{T_j} \right\rceil \text{len}(s_{max}^j(\Theta_j^h)) \right)}{T_i} \end{aligned} \quad (11)$$

Since FBLT is required to bound the effect of transitive retry, only  $\theta_i$  (not the whole  $\theta_i^{ex}$ ) will be considered in (11). Thus, ECM acts as if there were no transitive retry. Consequently, (11) holds if:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{s_i^k \in s_i} \delta_i^k \text{len}(s_i^k)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\sum_{\tau_j \in \gamma_i} \sum_{s_j^h(\Theta), \Theta \in (\theta_i \cap \theta_j^h)} \left( \left\lceil \frac{T_i}{T_j} \right\rceil \text{len}(s_{max}^j(\Theta)) \right)}{T_i} \end{aligned} \quad (12)$$

where  $s_{max}^j(\Theta) \leq s_{max}^j(\Theta_j^h)$ .

For each  $s_i^k \in s_i$ , there are a set of zero or more  $s_j^h(\Theta) \in \tau_j, \forall \tau_j \neq \tau_i$  that are conflicting with  $s_i^k$ . Assuming this set of transactions conflicting with  $s_i^k$  is denoted as  $\nu_i^k = \left\{ s_j^h(\Theta) \in \tau_j : (\Theta \in \theta_i \cap \theta_j^h) \wedge (\forall \tau_j \neq \tau_i) \wedge (s_j^h(\Theta) \notin \nu_i^l, l \neq k) \right\}$ .

The last condition  $s_j^h(\Theta) \notin \nu_i^l, l \neq k$  in the definition of  $\nu_i^k$  ensures that common transactions  $s_j^h$  that can conflict with more than one transaction  $s_i^k \in \tau_i$  are split among different  $\nu_i^k, k = 1, \dots, |s_i|$ . This condition is necessary, because in ECM, no two or more transactions of  $\tau_i^x$  can be aborted by the same transaction of  $\tau_j^h$ , where  $p_j^h > p_i^x$ . By substitution of  $\nu_i^k$  in (12), we get:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{s_i^k \in s_i} \delta_i^k \text{len}(s_i^k)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \sum_{k=1}^{|s_i|} \sum_{s_j^h(\Theta) \in \nu_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil \text{len}(s_{max}^j(\Theta)) \right) \right)}{T_i} \end{aligned} \quad (13)$$

(13) holds if for each  $s_i^k \in \tau_i$ :

$$\delta_i^k \leq \frac{\sum_{s_j^h(\Theta) \in \nu_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil \text{len}(s_{max}^j(\Theta)) \right)}{\text{len}(s_i^k)} \quad (14)$$

Since  $\text{len}(s_{max}^j(\Theta)) \geq \text{len}(s_i^k)$ , (14) holds if  $\delta_i^k \leq \sum_{s_j^h(\Theta) \in \nu_i^k} \left\lceil \frac{T_i}{T_j} \right\rceil \cdot \sum_{s_j^h(\Theta) \in \nu_i^k} \left\lceil \frac{T_i}{T_j} \right\rceil$  is the maximum number of transactions directly conflicting with  $s_i^k$  in all jobs with higher priority than  $\tau_i$ . Claim follows.  $\square$

## 7.2. FBLT vs. RCM

CLAIM 8. *The schedulability of FBLT is equal to or better than RCM's if*

$$\delta_i^k \leq \left( \sum_{s_j^h(\Theta) \in \nu_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \right) - \sum_{u=1, s_{u_{max}} \in \epsilon}^{\min(n,m)-1} s_{u_{max}}$$

$\sum_{s_j^h(\Theta) \in \nu_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right)$  is number of transactions directly conflicting with  $s_i^k$  in all jobs with higher priority than  $\tau_i$ .  $\sum_{u=1, s_{u_{max}} \in \epsilon}^{\min(n,m)-1} s_{u_{max}}$  is the sum of the maximum  $m - 1$  transactional lengths in all tasks

PROOF. By substituting  $RC_A(T_i)$  and  $RC_B(T_i)$  in (8) with (5) and (6.9) in [El-Shambakey 2012], respectively, we get:

$$\sum_{\forall \tau_i} \frac{\sum_{s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right) + RC_{re}(T_i)}{T_i} \quad (15)$$

$$\leq \sum_{\forall \tau_i} \frac{\left( \sum_{\forall \tau_j^* \in \gamma_i^{ex}} \sum_{\forall \theta \in \theta_i^{ex}} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \sum_{\forall s_j^h(\theta)} \text{len}(s_j^h(\theta) + s_{max}^j(\theta)) \right) + RC_{re}(T_i)}{T_i}$$

where  $\tau_j^* = \{\tau_j : (\tau_j \neq \tau_i) \wedge (p_j > p_i)\}$ .

Let  $\theta_i^{ex} = \theta_i + \theta_i^*$ , where  $\theta_i^*$  is the set of objects not directly accessed by any job of  $\tau_i$ , but can cause transactions in  $\tau_i$  to retry due to transitive retry. Let  $\gamma_i^{ex} = \gamma_i + \gamma_i^*$ , where  $\gamma_i^*$  is the set of tasks that access objects in  $\theta_i^*$ .  $s_j^h(\theta)$  can access multiple objects, so  $s_{max}^j(\theta)$  is the maximum length transaction conflicting with  $s_j^h(\theta)$ .  $s_j^h(\theta)$  is included only once for all  $\theta \in \Theta_j^h$ . Each  $\theta \in \Theta_j^h$  has its own  $s_{max}^j(\theta)$ . But  $s_i^h$  can access multiple objects, denoted as  $\Theta_j^h$ . So,  $s_{max}^j(\theta)$  is replaced by  $s_{max}^j(\Theta_j^h)$ , where  $s_{max}^j(\Theta_j^h) = \max\{s_{max}^j(\theta), \forall \theta \in \Theta_j^h\}$ .  $s_{max}^j(\Theta_j^h)$  is included once for each  $\theta \in \theta_i$ .

Each  $\tau_i^x$  has the same interference pattern from higher priority jobs,  $\tau_j^h$ , under FBLT and RCM. Hence,  $RC_{re}(T_i)$  for  $\tau_i^x$  is the same under FBLT and RCM. Consequently, (15) becomes:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\sum_{\tau_j^* \in \gamma_i^{ex}} \sum_{\forall s_j^h(\Theta_j^h), \Theta_j^h \in \theta_i^{ex}} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \text{len}(s_j^h(\Theta_j^h) + s_{max}^j(\Theta_j^h))}{T_i} \end{aligned} \quad (16)$$

Although different  $s_i^k$ s can have common conflicting transactions  $s_j^h$ , no more than one  $s_i^k$  can be preceded by the same  $s_j^h$  in the  $m\_set$ . This happens because transactions in the  $m\_set$  are non-preemptive. The original priority of transactions preceding  $s_i^k$  in the  $m\_set$  can be of lower or higher priority than the original priority of  $s_i^k$ . Under G-RMA,  $p_j > p_i$ , which means that  $T_j \leq T_i$ . Therefore,  $\left\lceil \frac{T_i}{T_j} \right\rceil \geq 1$ . For each  $s_i^k \in s_i$ , there are a set of zero or more  $s_j^h(\Theta_j^h) \in \tau_j^*$  that are conflicting with  $s_i^k$ . Assuming this set of transactions conflicting with  $s_i^k$  is denoted as  $\nu_i^k = \{s_j^h(\Theta_j^h) \in \tau_j^* : (\Theta_j^h \in \theta_i^{ex}) \wedge (s_j^h(\Theta_j^h) \notin \nu_i^l, l \neq k)\}$ .

The last condition  $s_j^h(\theta) \notin \nu_i^l, l \neq k$  in the definition of  $\nu_i^k$  ensures that common transactions  $s_j^h$  that can conflict with more than one transaction  $s_i^k \in \tau_i$  are split among different  $\nu_i^k, k = 1, \dots, |s_i|$ . This condition is necessary, because in RCM, no two or more transactions of  $\tau_i^x$  can be aborted by the same transaction of  $\tau_j^h$ , where  $p_j^h > p_i^x$ . By substitution of  $\nu_i^k$  in (16), we get:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \sum_{k=1}^{|s_i|} \sum_{s_j^h(\Theta_j^h) \in \nu_i^k} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \text{len}(s_j^h(\Theta_j^h) + s_{max}^j(\Theta_j^h)) \right) \right)}{T_i} \end{aligned} \quad (17)$$

$s_j^h$  belongs to higher priority jobs than  $\tau_i$ .  $s_{max}^j$  belongs to higher priority jobs than  $\tau_i$  or  $\tau_i$  itself.  $s_{max}^j$  has a lower priority than  $\tau_j$ . Transactions in the  $m\_set$  can belong to jobs with original priority higher or lower than  $\tau_i$ . Thus, (17) holds if for each  $s_i^k \in \tau_i$ :

$$\delta_i^k \text{len}(s_i^k) \leq \left( \sum_{s_j^h(\Theta_j^h) \in \nu_i^k} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \text{len}(s_j^h(\Theta_j^h) + s_{max}^j(\Theta_j^h)) \right) \right) - \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \quad (18)$$

Then,

$$\delta_i^k \leq \left( \sum_{\bar{s}_j^h(\Theta_j^h) \in \nu_i^k} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \text{len} \left( \frac{\bar{s}_j^h(\Theta_j^h) + s_{max}^j(\Theta_j^h)}{s_i^k} \right) \right) \right) - \sum_{s_{iz}^k \in \chi_i^k} \text{len} \left( \frac{s_{iz}^k}{s_i^k} \right) \quad (19)$$

Let  $\epsilon = \{s_{u_{max}} : (1 \leq u \leq n) \wedge (s_{u1_{max}} \geq s_{u2_{max}}, u1 < u2)\}$ , where  $n$  is the number of tasks, and  $s_{u_{max}}$  is the maximum transactional length in any job of  $\tau_u$ . Thus,  $\epsilon$  is the set of maximum transactional lengths of all tasks in non-increasing order. Each  $s_{u_{max}} \in \epsilon$  belongs to a distinct task. Thus,  $\sum_{s_{iz}^k \in \chi_i^k} \text{len} \left( \frac{s_{iz}^k}{s_i^k} \right) \leq \sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}}$ .  $\sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}}$  is the sum of at most maximum  $m - 1$  transactional lengths of all tasks.  $|\chi_i^k| \leq m - 1$  and  $\text{len}(s_{max}^j(\Theta_j^h)) \geq \text{len}(s_i^k)$ . So, (19) holds if:

$$\delta_i^k \leq \left( \sum_{\bar{s}_j^h(\Theta_j^h) \in \nu_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \right) - \sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}} \quad (20)$$

To bound the effect of transitive retry, only objects that belong to  $\theta_i$  (not whole  $\theta_i^{ex}$ ) will be considered. Thus, RCM acts as if there were no transitive retry. Thus,  $\nu_i^k$  is modified to  $\bar{\nu}_i^k = \left\{ \bar{s}_j^h(\Theta) \in \tau_j^* : (\Theta \in \Theta_j^h \cap \theta_i) \wedge (\bar{s}_j^h(\Theta) \notin \nu_i^l, l \neq k) \right\}$ . Since  $\bar{\nu}_i^k \subseteq \nu_i^k$ , (20) still holds if  $\nu_i^k$  is replaced with  $\bar{\nu}_i^k$ . Consequently, (20) holds if:

$$\delta_i^k \leq \left( \sum_{\bar{s}_j^h(\Theta) \in \bar{\nu}_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \right) - \sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}} \quad (21)$$

$\sum_{\bar{s}_j^h(\Theta) \in \bar{\nu}_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right)$  represents the number of transactions directly conflicting with  $s_i^k$  in all jobs with higher priority than  $\tau_i$ . Claim follows.  $\square$

### 7.3. FBLT vs. G-EDF/LCM

CLAIM 9. *The schedulability of FBLT is equal to or better than G-EDF/LCM's when*

$$\delta_i^k \leq \left( \sum_{\bar{s}_j^h(\Theta) \in \nu_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil \alpha_{max}^{j\bar{h}} \right) \right)$$

$\alpha_{max}^{j\bar{h}}$  is the maximum  $\alpha$  with which  $s_j^h$  can conflict with the maximum length transaction sharing objects with  $s_i^k$  and  $s_j^h$

PROOF. 7. The proof is similar to that of Claim 7 and is therefore skipped for brevity. It can be found in [El-Shambakey and Ravindran 2011].  $\square$

### 7.4. FBLT vs. G-RMA/LCM

CLAIM 10. *The schedulability of FBLT is equal to or better than G-RMA/LCM's when*

$$\delta_i^k \leq \left( \sum_{\bar{s}_j^h(\Theta) \in \bar{\nu}_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \alpha_{max}^{j\bar{h}} \right) - \sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}}$$



$\left(\sum_{s_j^h(\Theta) \in \nu_i^k} \left(\left\lceil \frac{T_i}{T_j} \right\rceil + 1\right)\right)$  is the sum of the total number each transaction  $s_j^h$  can directly conflict with  $s_i^k$ .  $\alpha_{max}^{\bar{j}^h}$  is the maximum  $\alpha$  with which  $s_j^h$  can conflict with the maximum length transaction sharing objects with  $s_i^k$  and  $s_j^h$ .

PROOF. The proof is similar to that of Claim 8 and is therefore skipped for brevity. It can be found in [El-Shambakey and Ravindran 2011].

□

## 7.5. FBLT vs. PNF

CLAIM 11. Let  $\rho_i^j(k) = \left(\sum_{\forall s_j^h(\Theta) \in \nu_i^k(j)} \text{len}(s_j^h(\Theta))\right) - s_{i_{max}}$ ,  $\tau_j \in \gamma_i^k$ .  $\rho_i^j(k)$  is the difference between the sum of transactional lengths of all transactions in  $\tau_j$  conflicting with  $s_i^k$ , and the maximum length transaction in  $\tau_i$ . Let  $\sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}}$  be sum of at most maximum  $m - 1$  transactional lengths in all tasks. Schedulability of FBLT is better or equal to PNF's when

$$\delta_i^k \leq \left(\sum_{\forall \tau_j \in \gamma_i^k} \left(\left\lceil \frac{T_i}{T_j} \right\rceil + 1\right) \text{len} \left( \frac{\left(\sum_{\forall s_j^h(\Theta) \in \nu_i^k(j)} \text{len}(s_j^h(\Theta))\right) - s_{i_{max}}}{s_i^k} \right)\right) - \sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}}$$

PROOF. By substituting  $RC_A(T_i)$  and  $RC_B(T_i)$  in (8) with (5) and (6.1) in [El-Shambakey 2012], respectively, we get:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{\forall s_i^k \in s_i} \left(\delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k)\right) + RC_{re}(T_i)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\sum_{\forall \tau_j \in \gamma_i} \sum_{\theta \in \theta_i} \left(\left(\left\lceil \frac{T_i}{T_j} \right\rceil + 1\right) \sum_{\forall s_j^h(\theta)} \text{len}(s_j^h(\theta))\right)}{T_i} \end{aligned} \quad (22)$$

$s_j^h(\theta)$  can access multiple objects.  $s_j^h(\theta)$  is included only once for all objects accessed by it.  $RC_{re}(T_i)$  is given by (6.8) in [El-Shambakey 2012] in case of G-EDF, and (6.10) in [El-Shambakey 2012] in case of G-RMA. Substituting  $RC_{re}(T_i)$  given by (6.8) and (6.10) in [El-Shambakey 2012] with  $RC_{re}(T_i) = \sum_{\forall \tau_j \in \gamma_i} \left(\left\lceil \frac{T_i}{T_j} \right\rceil + 1\right) s_{i_{max}}$ , we ensure correctness of (22). If  $\tau_j$  has no shared objects with  $\tau_i$ , then the release of any higher priority job  $\tau_j^y$  will not abort any transaction in any job of  $\tau_i$ . Thus, (22) holds if:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{\forall s_i^k \in s_i} \left(\delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k)\right)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\sum_{\forall \tau_j \in \gamma_i} \left(\left\lceil \frac{T_i}{T_j} \right\rceil + 1\right) \left(\left(\sum_{\forall s_j^h(\Theta), \Theta \in \theta_i} \text{len}(s_j^h(\Theta))\right) - s_{i_{max}}\right)}{T_i} \end{aligned} \quad (23)$$

For each  $s_i^k \in s_i$ , there are a set of zero or more  $s_j^h(\Theta) \in \tau_j, \forall \tau_j \neq \tau_i$  that are conflicting with  $s_i^k$ . Assuming this set of transactions conflicting with  $s_i^k$  is denoted as  $\nu_i^k(j) = \left\{s_j^h(\Theta) \in \tau_j : (\Theta \in \theta_i) \wedge (\tau_j \neq \tau_i) \wedge \left(s_j^h(\Theta) \notin \nu_i^l, l \neq k\right)\right\}$ . The last condition  $s_j^h(\Theta) \notin \nu_i^l, l \neq k$  in the definition of  $\nu_i^k$  ensures that common transactions  $s_j^h$  that can conflict with more than one transaction  $s_i^k \in \tau_i$  are split among different  $\nu_i^k(j), k = 1, \dots, |s_i|$ . This condition is necessary, because in PNF, no two or more transactions of  $\tau_i^x$  can be aborted by the same transaction of  $\tau_j^h$ .

Let  $\gamma_i^k$  be the subset of  $\gamma_i$  that contains tasks with transactions conflicting directly with  $s_i^k$ . By substitution of  $\nu_i^k$  and  $\gamma_i$  in (23) by  $\nu_i^k(j)$  and  $\gamma_i^k$ , (23) holds if for each  $s_i^k$ :

$$\begin{aligned} & \delta_i^k + \sum_{s_{iz}^k \in \chi_i^k} \text{len}\left(\frac{s_{iz}^k}{s_i^k}\right) \\ & \leq \sum_{\forall \tau_j \in \gamma_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \left( \left( \sum_{\forall s_j^h(\Theta) \in \nu_i^k(j)} \text{len}\left(\frac{s_j^h(\Theta)}{s_i^k}\right) \right) - \text{len}\left(\frac{s_{imax}}{s_i^k}\right) \right) \end{aligned} \quad (24)$$

Non-preemptive transactions preceding  $s_i^k$  in the  $m$ -set can directly or indirectly conflict with  $s_i^k$ . Under PNF, transactions can only directly conflict with  $s_i^k$ . Thus,  $s_{iz}^k$  on the left hand side of (24) is not necessarily included in  $\bar{s}_j^h(\Theta)$  on the right hand side of (24). Let  $\epsilon = \{s_{u_{max}} : (1 \leq u \leq n) \wedge (s_{u1_{max}} \geq s_{u2_{max}}, u1 < u2)\}$ , where  $n$  is the number of tasks, and  $s_{u_{max}}$  is the maximum transactional length in any job of  $\tau_u$ . Thus,  $\epsilon$  is the set of maximum transactional lengths of all tasks in non-increasing order. Each  $s_{u_{max}} \in \epsilon$  belongs to a distinct task. Thus,  $\sum_{s_{iz}^k \in \chi_i^k} \text{len}\left(\frac{s_{iz}^k}{s_i^k}\right) \leq \sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}}$ .  $\sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}}$  is the sum of at most maximum  $m-1$  transactional lengths of all tasks.  $|\chi_i^k| \leq m-1$ . Then (24) holds if:

$$\begin{aligned} \delta_i^k & \leq \left( \sum_{\forall \tau_j \in \gamma_i^k} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \text{len}\left(\frac{\left( \sum_{\forall s_j^h(\Theta) \in \nu_i^k(j)} \text{len}\left(\bar{s}_j^h(\Theta)\right) \right) - s_{imax}}{s_i^k}\right) \right) \\ & \quad - \sum_{u=1, s_{u_{max}} \in \epsilon}^{min(n,m)-1} s_{u_{max}} \end{aligned}$$

Since  $\rho_i^j(k) = \left( \sum_{\forall s_j^h(\Theta) \in \nu_i^k(j)} \text{len}\left(\bar{s}_j^h(\Theta)\right) \right) - s_{imax}$ ,  $\tau_j \in \gamma_i^k$ , and  $\left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right)$  is the total number of jobs of  $\tau_j$  interfering with  $\tau_i$ , claim follows.  $\square$

## 7.6. FBLT vs. Lock-free

**CLAIM 12.** *Under G-EDF and G-RMA, the schedulability of FBLT is equal or better than that under lock-free synchronization if  $s_{max} \leq r_{max}$ . If transactions execute in FIFO order (i.e.,  $\delta_i^k = 0, \forall s_i^k$ ) and contention is high,  $s_{max}$  can be much larger than  $r_{max}$ .*

**PROOF.** Lock-free synchronization [Devi et al. 2006; Herlihy 2006] allows only one object to be synchronized at a given time (e.g., a lock-free stack). Thus, for comparing FBLT's schedulability with lock-free, we limit the number of accessed objects per transaction under FBLT to one.

By substituting  $RC_A(T_i)$  and  $RC_B(T_i)$  in (8) with (5) and (6.17) in [El-Shambakey 2012], respectively, we get:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right) + RC_{re}(T_i)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} r_{max} \right) \right) + RC_{re}(T_i)}{T_i} \end{aligned} \quad (25)$$

where  $\beta_{i,j}$  is the number of retry loops of  $\tau_j$  that access the same objects as accessed by any retry loop of  $\tau_i$  [Devi et al. 2006] and  $r_{max}$  is the maximum execution cost of a single iteration of any retry loop of any task [Devi et al. 2006].

For G-EDF (G-RMA), any job  $\tau_i^x$  under FBLT has the same pattern of interference from higher priority jobs as ECM (RCM), respectively.  $RC_{re}(T_i)$  for ECM, RCM, and lock-free are given by Claims 25, 26, and 27 in [El-Shambakey 2012], respectively.  $RC_{re}(T_i) = \left\lceil \frac{T_i}{T_j} \right\rceil s_{imax}, \forall \tau_j \neq \tau_i$  for G-EDF/FBLT and G-RMA/FBLT.  $RC_{re}(T_i) = \left\lceil \frac{T_i}{T_j} \right\rceil r_{imax}, \forall \tau_j \neq \tau_i$  for G-EDF/lock-free and G-RMA/lock-free. (31) becomes:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{\forall s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil s_{imax}}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} r_{max} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil r_{imax}}{T_i} \end{aligned} \quad (26)$$

Since  $s_{max} \geq s_{imax}$ ,  $\text{len}(s_i^k)$ ,  $\text{len}(s_{iz}^k)$ ,  $\forall i, z, k$  and  $r_{max} \geq r_{imax}$  (32) holds if:

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\left( \left( \sum_{\forall s_i^k \in s_i} (\delta_i^k + |\chi_i^k|) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil \right) s_{max}}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil \right) r_{max}}{T_i} \end{aligned} \quad (27)$$

(33) holds if for each  $\tau_i$ :

$$\begin{aligned} & \left( \left( \sum_{\forall s_i^k \in s_i} (\delta_i^k + |\chi_i^k|) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil \right) s_{max} \\ & \leq \left( \left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil \right) r_{max} \end{aligned} \quad (28)$$

$$\frac{s_{max}}{r_{max}} \leq \frac{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil}{\left( \sum_{\forall s_i^k \in s_i} (\delta_i^k + |\chi_i^k|) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil} \quad (29)$$

It appears from (35) that as  $\delta_i^k$  and  $|\chi_i^k|$  increases,  $s_{max}/r_{max}$  decreases. So, to get the lower bound on  $s_{max}/r_{max}$ , let  $\sum_{\forall s_i^k \in s_i} (\delta_i^k + |\chi_i^k|)$  reach its maximum value. This maximum value is the total number of interfering transactions belonging to any job  $\tau_j^l$ ,  $j \neq i$ . The priority of  $\tau_j^l$  can be higher or lower than the current instance of  $\tau_i$ . Beyond this maximum value, there will be no more transactions that conflict with  $s_i^k$ . Thus, higher values for any  $\delta_i^k$  beyond the maximum value will be ineffective.  $\sum_{\forall s_i^k \in s_i} (\delta_i^k + |\chi_i^k|) \leq \sum_{\forall \tau_j \in \gamma_i} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right)$ . Consequently, (35) will be:

$$\frac{s_{max}}{r_{max}} \leq \frac{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil}{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil} \quad (30)$$

Since we are seeking the lower bound on  $\frac{s_{max}}{r_{max}}$ , let  $\beta_{i,j}$  assume its minimum value. Thus,  $\beta_{i,j} = 1$ . (36) holds if  $\frac{s_{max}}{r_{max}} \leq 1$ .

Let  $\delta_i^k(T_i) \rightarrow 0$  in (35). This means that transactions approximately execute according to their arrival order. Let  $\beta_{i,j} \rightarrow \infty$ ,  $\left\lceil \frac{T_i}{T_j} \right\rceil \rightarrow \infty$  in (35). This implies high contention. Consequently,  $\frac{s_{max}}{r_{max}} \rightarrow \infty$ . Hence, if transactions execute in FIFO order and contention is high,  $s_{max}$  can be much larger than  $r_{max}$ . Claim follows.  $\square$

## CLAIM 13.

Under G-EDF and G-RMA, schedulability of FBLT is equal or better than lock-free's if  $s_{max} \leq r_{max}$ . If transactions execute in FIFO order (i. e.,  $\delta_i^k = 0, \forall s_i^k$ ) and contention is high,  $s_{max}$  can be much larger than  $r_{max}$ .

## PROOF.

Lock-free synchronization [Devi et al. 2006; El-Shambakey and Ravindran 2012b] accesses only one object. Thus, the number of accessed objects per transaction in FBLT is limited to one. This allows us to compare the schedulability of FBLT with the lock-free algorithm.

By substituting  $RC_A(T_i)$  and  $RC_B(T_i)$  in (8) with (5) and (6.17) in [El-Shambakey 2012] respectively.

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right) + RC_{re}(T_i)}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} r_{max} \right) \right) + RC_{re}(T_i)}{T_i} \end{aligned} \quad (31)$$

where  $\beta_{i,j}$  is the number of retry loops of  $\tau_j$  that access the same objects as accessed by any retry loop of  $\tau_i$  [Devi et al. 2006].  $r_{max}$  is the maximum execution cost of a single iteration of any retry loop of any task [Devi et al. 2006]. For G-EDF(G-RMA), any job  $\tau_i^x$  under FBLT has the same pattern of interference from higher priority jobs as ECM(RCM) respectively.  $RC_{re}(T_i)$  for ECM, RCM and lock-free are given by Claims 25, 26 and 27 in [El-Shambakey 2012] respectively.  $RC_{re}(T_i) = \left\lceil \frac{T_i}{T_j} \right\rceil s_{imax}, \forall \tau_j \neq \tau_i$

covers  $RC_{re}(T_i)$  for G-EDF/FBLT and G-RMA/FBLT.  $RC_{re}(T_i) = \left\lceil \frac{T_i}{T_j} \right\rceil r_{imax}, \forall \tau_j \neq \tau_i$  covers retry cost for G-EDF/lock-free and G-RMA/lock-free. (31) becomes

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\sum_{s_i^k \in s_i} \left( \delta_i^k \text{len}(s_i^k) + \sum_{s_{iz}^k \in \chi_i^k} \text{len}(s_{iz}^k) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil s_{imax}}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} r_{max} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil r_{imax}}{T_i} \end{aligned} \quad (32)$$

Since  $s_{max} \geq s_{imax}$ ,  $\text{len}(s_i^k), \text{len}(s_{iz}^k), \forall i, z, k$  and  $r_{max} \geq r_{imax}$  (32) holds if

$$\begin{aligned} & \sum_{\forall \tau_i} \frac{\left( \left( \sum_{s_i^k \in s_i} (\delta_i^k + |\chi_i^k|) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil \right) s_{max}}{T_i} \\ & \leq \sum_{\forall \tau_i} \frac{\left( \left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil \right) r_{max}}{T_i} \end{aligned} \quad (33)$$

(33) holds if for each  $\tau_i$

$$\begin{aligned} & \left( \left( \sum_{s_i^k \in s_i} (\delta_i^k + |\chi_i^k|) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil \right) s_{max} \\ & \leq \left( \left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil \right) r_{max} \end{aligned} \quad (34)$$

$$\frac{s_{max}}{r_{max}} \leq \frac{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil}{\left( \sum_{s_i^k \in s_i} (\delta_i^k + |\chi_i^k|) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil} \quad (35)$$

It appears from (35) that as  $\delta_i^k$ , as well as  $|\chi_i^k|$ , increases, then  $s_{max}/r_{max}$  decreases. So, to get the lower bound on  $s_{max}/r_{max}$ , let  $\sum_{\forall s_i^k \in s_i} (\delta_i^k + |\chi_i^k|)$  reaches its maximum value. This maximum value is the total number of interfering transactions belonging to any job  $\tau_j^l$ ,  $j \neq i$ . Priority of  $\tau_j^l$  can be higher or lower than current instance of  $\tau_i$ . Beyond this maximum value, there will be no more transactions to conflict with  $s_i^k$ . So, higher values for any  $\delta_i^k$  beyond maximum value will be ineffective.  $\sum_{\forall s_i^k \in s_i} (\delta_i^k + |\chi_i^k|) \leq \sum_{\forall \tau_j \in \gamma_i} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right)$ . Consequently, (35) will be

$$\frac{s_{max}}{r_{max}} \leq \frac{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \beta_{i,j} \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil}{\left( \sum_{\forall \tau_j \in \gamma_i} \left( \left\lceil \frac{T_i}{T_j} \right\rceil + 1 \right) \right) + \sum_{\forall \tau_j \neq \tau_i} \left\lceil \frac{T_i}{T_j} \right\rceil} \quad (36)$$

As we look for the lower bound on  $\frac{s_{max}}{r_{max}}$ , let  $\beta_{i,j}$  assumes its minimum value. So,  $\beta_{i,j} = 1$ . (36) holds if  $\frac{s_{max}}{r_{max}} \leq 1$ .

Let  $\delta_i^k(T_i) \rightarrow 0$  in (35). This means transactions approximately execute in their arrival order. Let  $\beta_{i,j} \rightarrow \infty$ ,  $\left\lceil \frac{T_i}{T_j} \right\rceil \rightarrow \infty$  in (35). This means contention is high. Consequently,  $\frac{s_{max}}{r_{max}} \rightarrow \infty$ . So, if transactions execute in FIFO order and contention is high,  $s_{max}$  can be much larger than  $r_{max}$ . Claim follows.

□

## 8. EXPERIMENTAL EVALUATION

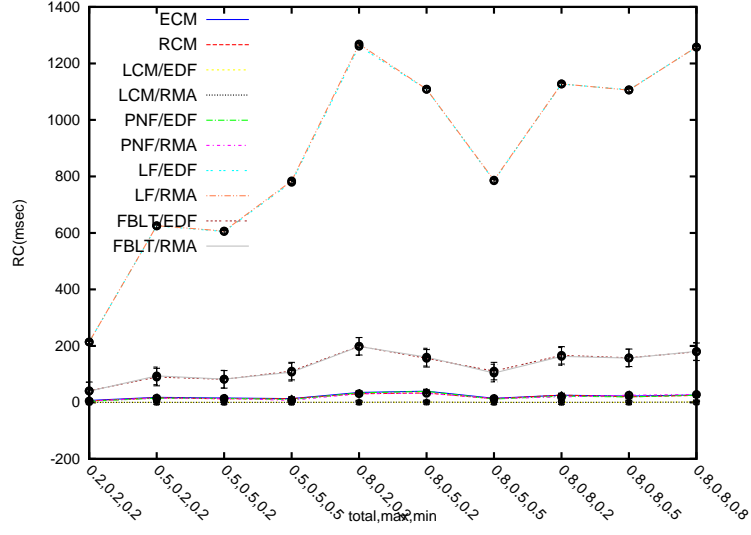
We now would like to understand how FBLT's retry cost compares with competitors in practice (i.e., on average). Since this can only be understood experimentally, we implement FBLT and the competitors and conduct experiments.

We used the ChronOS real-time Linux kernel [Dellinger et al. 2011] and the RSTM library [Marathe et al. ] in our implementation. We implemented G-EDF and G-RMA schedulers in ChronOS, and modified RSTM to include implementations of ECM, RCM, LCM, PNF and FBLT. For the retry-loop lock-free synchronization, we used a loop that reads an object and attempts to write to it using a CAS instruction. The task retries until the CAS succeeds. We used an 8 core, 2GHz AMD Opteron platform. The average time taken for one write operation by RSTM on any core is  $0.0129653375 \mu s$ , and the average time taken by one CAS-loop operation on any core is  $0.0292546250 \mu s$ .


We used four task sets consisting of 4, 5, 8, and 20 periodic tasks. Each task runs in its own thread and has a set of atomic sections. Atomic section properties are probabilistically controlled using three parameters: the maximum and minimum lengths of any atomic section within a task, and the total length of atomic sections within any task. Since lock-free synchronization cannot handle more than one object per atomic section, we first compare FBLT's retry cost with that of lock-free (and other CMs) for one object per transaction. We then compare FBLT's retry cost with that of other CMs for multiple objects per transaction.

Figures 1 show the average retry cost for the 5 task case sharing 1 object. On the x-axis of the figures, we record 3 parameters  $x$ ,  $y$ , and  $z$ .  $x$  is the ratio of the total length of all atomic sections of a task to the task WCET.  $y$  is the ratio of the maximum length of any atomic section of a task to the task WCET.  $z$  is the ratio of the minimum length of any atomic section of a task to the task WCET. Confidence level of all data points is 0.95.

While Figure 1(a) includes all methods, Figure 1(b) excludes lock-free. From these figures, we observe that lock-free has the largest retry cost, as it provides no conflict resolution. FBLT has the longest retry cost among CMs because transactions share



(a) ECM, RCM, LCM, PNF, Lock-Free



(b) ECM, RCM, LCM, PNF

Fig. 1. Avg. retry cost (one object/transaction).

only 1 object. For multi-objects per transaction, FBLT provides equal or shorter retry cost than LCM as shown in Figures 2 and 3. PNF has an advantage over FBLT. PNF knows a priori all objects accessed by each transaction, but FBLT does not. Consequently, retry cost under PNF is shorter than FBLT. Experiments show that by proper adjustment of maximum abort number of each transaction, retry cost can be shorter than ECM, RCM and LCM. Retry cost can be comparable to PNF.

Similar trends were observed for the other task sets those are omitted here for brevity and due to space limitations. Other results are shown in [El-Shambaakey and Ravindran 2011].




Fig. 2. Avg. retry cost (20 shared objects, 4 tasks).




Fig. 3. Avg. retry cost (40 shared objects, 4 tasks).

## 9. CONCLUSIONS

Transitive retry increases transactional retry cost under ECM, RCM, and LCM. PNF avoids transitive retry by avoiding transactional preemptions. PNF avoids retry cost by concurrently executing non-conflicting. Non-conflicting objects are non-preemptive. PNF requires priori knowledge about objects accessed by each transaction. This is not suitable with dynamic STM. So, we introduce the First Bounded, Last Timestamp (FBLT) CM. Each transaction is allowed to abort for a specified number. Afterwards, transaction becomes non-preemptive. Non-preemptive transactions have higher prior-

ities than other preemptive transactions and real-time jobs. Non-preemptive transactions resolve their conflicts using FIFO order.

By proper adjustment of maximum abort number of each transaction, FBLT schedulability is equal or better than other synchronization techniques. For FBLT's schedulability to be equal or better than lock-free, the upper bound on  $s_{max}/r_{max}$  is 1. Upper bound on  $s_{max}/r_{max}$  can be higher than 1 if transactions execute in their arrival order and contention is high.

Experimental results show that FBLT has equal or shorter retry cost than ECM, RCM and LCM. PNF knows a priori all objects accessed by each transaction. This is an advantage for PNF over FBLT. Consequently, retry cost under PNF is shorter than retry cost under FBLT.

Future work includes choosing another criterion to resolve conflicts of non-preemptive transactions. Also, using feedback from the system to adjust maximum abort number of each transaction. Consequently, retry cost can be reduced over time.

## REFERENCES

- ANDERSON, J., RAMAMURTHY, S., AND JEFFAY, K. 1995. Real-time computing with lock-free shared objects. In *RTSS*. 28–37.
- BARROS, A. AND PINHO, L. 2011. Managing contention of software transactional memory in real-time systems. In *IEEE RTSS, Work-In-Progress*.
- BRANDENBURG, B., CALANDRINO, J., BLOCK, A., LEONTYEV, H., AND ANDERSON, J. 2008. Real-time synchronization on multiprocessors: To block or not to block, to suspend or spin? In *Real-Time and Embedded Technology and Applications Symposium, 2008. RTAS '08. IEEE*. 342–353.
- DAVIS, R. I. AND BURNS, A. 2011. A survey of hard real-time scheduling for multiprocessor systems. *ACM Comput. Surv.* 43, 4, 35:1–35:44.
- DELLINGER, M., GARYALI, P., AND RAVINDRAN, B. 2011. Chronos linux: a best-effort real-time multiprocessor linux kernel. In *Proceedings of the 48th Design Automation Conference*. DAC '11. ACM, New York, NY, USA, 474–479.
- DEVI, U., LEONTYEV, H., AND ANDERSON, J. 2006. Efficient synchronization under global edf scheduling on multiprocessors. In *18th Euromicro Conference on Real-Time Systems*. 10 pp. –84.
- EL-SHAMBAKEY, M. 2012. Phd proposal. Ph.D. thesis, Virginia Tech.
- EL-SHAMBAKEY, M. AND RAVINDRAN, B. 2011. On the design of real-time stm contention managers. Tech. rep., ECE Department, Virginia Tech. Available as <http://www.real-time.ece.vt.edu/tech-report-rt-stm-cm11.pdf>.
- EL-SHAMBAKEY, M. AND RAVINDRAN, B. 2012a. Stm concurrency control for embedded real-time software with tighter time bounds. In *Proceedings of the 49th Annual Design Automation Conference*. DAC '12. ACM, New York, NY, USA, 437–446.
- EL-SHAMBAKEY, M. AND RAVINDRAN, B. 2012b. Stm concurrency control for multicore embedded real-time software: time bounds and tradeoffs. In *Proceedings of the 27th Annual ACM Symposium on Applied Computing*. SAC '12. ACM, New York, NY, USA, 1602–1609.
- FAHMY, S. AND RAVINDRAN, B. 2011. On stm concurrency control for multicore embedded real-time software. In *International Conference on Embedded Computer Systems*. SAMOS. 1–8.
- FAHMY, S., RAVINDRAN, B., AND JENSEN, E. D. 2009. On bounding response times under software transactional memory in distributed multiprocessor real-time systems. In *DATE*. 688–693.
- FAHMY, S. F. 2010. Collaborative scheduling and synchronization of distributable real-time threads. Ph.D. thesis, Virginia Tech.
- GUERRAoui, R., HERLIHY, M., AND POCHON, B. 2005. Toward a theory of transactional contention managers. In *PODC*. 258–264.
- HARRIS, T., MARLOW, S., JONES, S. P., AND HERLIHY, M. 2008. Composable memory transactions. *Commun. ACM* 51, 91–100.
- HERLIHY, M. 2006. The art of multiprocessor programming. In *PODC*. 1–2.
- HERLIHY, M., LUCHANGCO, V., MOIR, M., AND SCHERER, III, W. N. 2003. Software transactional memory for dynamic-sized data structures. In *Proceedings of the twenty-second annual symposium on Principles of distributed computing*. PODC '03. ACM, New York, NY, USA, 92–101.
- MANSON, J., BAKER, J., ET AL. 2006. Preemptible atomic regions for real-time Java. In *RTSS*. 10–71.



- MARATHE, V., SPEAR, M., HERIOT, C., ACHARYA, A., EISENSTAT, D., SCHERER III, W., AND SCOTT, M. Lowering the overhead of nonblocking software transactional memory. In *Workshop on Languages, Compilers, and Hardware Support for Transactional Computing*. TRANSACT.
- SAHA, B., ADL-TABATABAI, A.-R., ET AL. 2006. McRT-STM: a high performance software transactional memory system for a multi-core runtime. In *PPoPP*. 187–197.
- SARNI, T., QUEUDET, A., AND VALDURIEZ, P. 2009. Real-time support for software transactional memory. In *RTCSA*. 477–485.
- SCHOEBERL, M., BRANDNER, F., AND VITEK, J. 2010. RTTM: Real-time transactional memory. In *ACM SAC*. 326–333.

Received February 2007; revised March 2009; accepted June 2009